

# SmartExplain AI – Project Details

---

## Interpretable & Adaptive House Price Prediction Engine

---

## Table of Contents

1. [Project Overview](#)
  2. [How It Works \(End-to-End Flow\)](#)
  3. [Datasets](#)
  4. [Feature Engineering](#)
  5. [Model: LinearRegressionGD](#)
  6. [Optimization \(Gradient Descent\)](#)
  7. [Mathematical Formulation](#)
  8. [Explainability](#)
  9. [What-If Simulator](#)
  10. [Visualization](#)
  11. [Streamlit App](#)
  12. [Training Pipeline](#)
  13. [Project Structure](#)
  14. [Results](#)
- 

## 1. Project Overview

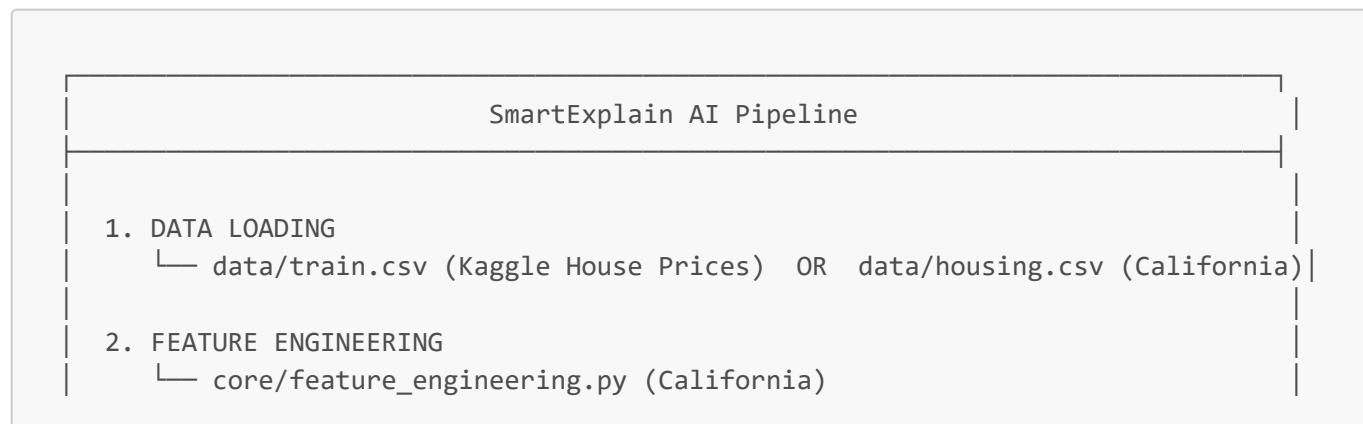
SmartExplain AI is a production-level machine learning system that predicts house prices using **linear regression trained via gradient descent** (not sklearn's closed-form solution). It emphasizes:

- **Interpretability** – per-feature contributions to each prediction
- **Adaptive optimization** – momentum, learning rate decay, early stopping
- **Modularity** – separate core, visualization, simulator, and app modules
- **Reproducibility** – fixed random seeds and clear pipeline

**Main model:** Custom `LinearRegressionGD` class. Sklearn's `LinearRegression` is used only for comparison.

---

## 2. How It Works (End-to-End Flow)



```
└── core/feature_engineering_house_prices.py (Kaggle)
    • Handle missing values
    • Outlier capping (1st-99th percentile)
    • Log transforms (skewed features)
    • Engineered features (interactions, location, etc.)
    • Polynomial expansion (squares + pairwise products)
    • Standardization:  $(x - \mu) / \sigma$ 

3. TRAIN/TEST SPLIT (80/20)
└── Random permutation, reproducible with seed

4. MODEL TRAINING
└── core/model.py - LinearRegressionGD
    • Batch gradient descent
    • Momentum ( $\beta=0.9$ )
    • Learning rate decay (time-based)
    • Early stopping (patience=150)
    • L2 regularization ( $\lambda=0.02$ )

5. RETRAIN ON FULL DATA
└── Final model trained on all samples for deployment

6. SAVE MODEL
└── model.pkl (weights, bias, feature engineer, feature names)

7. PREDICTION & EXPLAINABILITY
└── core/explainability.py - Contribution_i =  $w_i \times x_i$ 
└── simulator/what_if.py - simulate_price_change()

8. INTERACTIVE APP
└── app/streamlit_app.py - Sliders → prediction + contribution chart
```

### 3. Datasets

#### 3.1 Kaggle House Prices (Primary)

- **File:** `data/train.csv`
- **Source:** [House Prices - Advanced Regression Techniques](#)
- **Samples:** 1,460
- **Target:** `SalePrice` (house sale price in USD)
- **Features:** 80 columns (numeric + categorical)
- **Selected numeric:** LotArea, OverallQual, GrLivArea, GarageCars, TotalBsmtSF, YearBuilt, FullBath, Fireplaces, etc.

If `train.csv` exists in `data/`, this dataset is used.

#### 3.2 California Housing (Fallback)

- **File:** `data/housing.csv`
- **Samples:** ~20,000
- **Target:** `median_house_value`
- **Features:** longitude, latitude, housing\_median\_age, total\_rooms, total\_bedrooms, population, households, median\_income, ocean\_proximity

Used when `train.csv` is not present.

---

## 4. Feature Engineering

### 4.1 Kaggle House Prices (`HousePricesFeatureEngineer`)

1. **Numeric columns:** 34 columns (LotFrontage, LotArea, OverallQual, YearBuilt, GrLivArea, GarageCars, etc.)
2. **Missing values:** filled with median per column
3. **Standardization:**  $(x - \mu) / \sigma$
4. **Polynomial expansion (degree 2):**
  - Squares:  $x_i^2$  for first 8 features
  - Interactions:  $x_i \times x_j$  for pairs of first 8 features
5. **Output:** ~69 features

### 4.2 California Housing (`FeatureEngineer`)

1. **Base features:** longitude, latitude, median\_income, total\_rooms, total\_bedrooms, housing\_median\_age, population, households
  2. **Engineered features:**
    - `area_location_rating` = total\_rooms  $\times$  location\_rating (from ocean\_proximity)
    - `age_depreciation` =  $\exp(-0.02 \times age)$
    - `distance_from_center` =  $\sqrt{(\text{lon} - \text{lon}_c)^2 + (\text{lat} - \text{lat}_c)^2}$
    - `rooms_per_household, bedrooms_per_room`
    - `income_area, income_age, dist_location`
  3. **Log transforms:**  $\log_{10}(\text{total_rooms})$ ,  $\log_{10}(\text{bedrooms})$ ,  $\log_{10}(\text{population})$ , etc.
  4. **One-hot encoding:** ocean\_proximity
  5. **Outlier capping:** features clipped at 1st and 99th percentiles
  6. **Polynomial expansion:** degree 2 on first 10 features
  7. **Standardization:**  $(x - \mu) / \sigma$
- 

## 5. Model: LinearRegressionGD

**Location:** `core/model.py`

### 5.1 Formula

- **Prediction:**  $y_{\text{pred}} = Xw + b$
- **Cost:**  $J(w,b) = (1/2m) \sum (y_{\text{pred}} - y)^2 + \lambda \sum w^2$  (MSE + L2)
- **Gradients:**
  - $dw = (1/m) X^T(y_{\text{pred}} - y) + 2\lambda w$

- $db = (1/m) \sum(y_{pred} - y)$

## 5.2 Modes

Mode	Batch Size	Description
batch	m (full data)	Uses all samples per iteration
minibatch	32 (default)	Random subset per iteration
sgd	1	Single sample per iteration

## 5.3 Features

- **Momentum:**  $v = \beta \cdot v + \nabla J$ , update  $w$  using  $v$
- **Learning rate decay:**  $\alpha_t = \alpha_0 / (1 + \text{decay\_rate} \cdot t)$  (time decay)
- **Early stopping:** stop if cost does not improve for **patience** iterations
- **Gradient clipping:** prevents overflow during training

## 5.4 Stored Artifacts

- **weights** ( $w$ )
- **bias** ( $b$ )
- **cost\_history** (cost per iteration)

# 6. Optimization (Gradient Descent)

## 6.1 Momentum (`core/optimizers.py`)

$$v = \beta \times v + \nabla J$$

$$w = w - \alpha \times v$$

- $\beta = 0.9$
- Smoother updates and faster convergence in stable directions

## 6.2 Learning Rate Scheduler

Type	Formula
Time	$\alpha_t = \alpha_0 / (1 + \gamma \cdot t)$
Step	$\alpha_t = \alpha_0 \times \gamma^{(t // \text{step})}$
Exponential	$\alpha_t = \alpha_0 \times \exp(-\gamma \cdot t)$

Used: **time decay** with  $\gamma = 0.01$ .

# 7. Mathematical Formulation

## 7.1 Linear Model

$$\$y = Xw + b\$$$

## 7.2 Cost Function (L2 Regularized)

$$\$J(w,b) = \frac{1}{2m} \sum_{i=1}^m (y_{\text{pred}}^{(i)} - y^{(i)})^2 + \lambda \sum_j w_j^2 \$$$

## 7.3 Gradient Updates

$$\$\frac{\partial J}{\partial w} = \frac{1}{m} X^T (y_{\text{pred}} - y) + 2\lambda w \$$$

$$\$\frac{\partial J}{\partial b} = \frac{1}{m} \sum_i (y_{\text{pred}}^{(i)} - y^{(i)}) \$$$

## 7.4 Training Step

1. Compute  $y_{\text{pred}} = Xw + b$
  2. Compute  $dw, db$
  3. Apply momentum:  $v = \beta \cdot v + dw$
  4. Update:  $w \leftarrow w - \alpha \cdot v, b \leftarrow b - \alpha \cdot db$
- 

## 8. Explainability

**Location:** [core/explainability.py](#)

### 8.1 Feature Contribution

For linear model  $y = w \cdot x + b$ :

$$\$ \text{Contribution}_i = w_i \times x_i \$$$

### 8.2 Output

- **total\_prediction:** predicted price
- **contributions:** per-feature contribution ( $w_i \times x_i$ )
- **percentages:** relative influence =  $100 \times |\text{contribution}_i| / \sum |\text{contribution}_j|$

### 8.3 Use Case

Shows which features (e.g., GrLivArea, OverallQual) most influence each prediction.

---

## 9. What-If Simulator

**Location:** [simulator/what\\_if.py](#)

### 9.1 Function

```
simulate_price_change(model, X_current, feature_name, new_value, feature_names,
explainer)
```

## 9.2 Returns

- `original_prediction` – price before change
- `updated_prediction` – price after change
- `price_difference` –  $\Delta = \text{updated} - \text{original}$
- `contribution_breakdown_original` – contributions before
- `contribution_breakdown_updated` – contributions after

## 9.3 Example

Change `GrLivArea` from 1710 to 2500 → get new price and updated contribution breakdown.

---

# 10. Visualization

## 10.1 `visualization/plots.py`

- **Cost vs iterations** – training cost curve
- **Optimizer comparison** – batch vs minibatch vs SGD
- **Actual vs predicted** – scatter plot
- **Learning rate comparison** – convergence for different  $\alpha$

## 10.2 `visualization/cost_surface.py`

- **3D cost surface** –  $J(w_1, w_2)$  vs  $w_1, w_2$
  - **Gradient descent path** – trajectory over the surface
- 

# 11. Streamlit App

**Location:** `app/streamlit_app.py`

## 11.1 House Prices Mode (Kaggle)

Sliders:

- Lot Area, Overall Quality, Gr Liv Area, Garage Cars
- Total Basement Sq Ft, Year Built, Full Bath, Fireplaces

## 11.2 California Mode

Sliders:

- Area, Bedrooms, Location Rating, Age, Distance
- Median Income, Population, Households
- Longitude, Latitude, Ocean Proximity

## 11.3 Display

- Predicted price

- Feature contribution bar chart
  - Training cost curve
- 

## 12. Training Pipeline

### 12.1 train.py Flow

- Detect dataset:** use `train.csv` if present, else `housing.csv`
- Feature engineering:** call `fe.fit_transform(df)`
- Split:** 80% train, 20% test (random, seeded)
- Train:** `LinearRegressionGD.fit(X_train, y_train)`
- Retrain:** fit again on full data
- Evaluate:** MAE, RMSE,  $R^2$  on train, test, and full data
- Save:** pickle `{model, fe, feature_names, dataset}` to `model.pkl`

### 12.2 Hyperparameters (Current)

Parameter	Value
learning_rate	0.03
n_iterations	5000
regularization	0.02
momentum	0.9
decay_type	time
patience	150
random_state	42

---

## 13. Project Structure

```

SmartExplain-AI/
    ├── data/
    │   ├── housing.csv          # California housing
    │   ├── train.csv            # Kaggle House Prices
    │   ├── test.csv
    │   ├── data_description.txt
    │   └── house_prices.zip
    ├── notebooks/
    │   ├── SmartExplain_AI.ipynb
    │   └── SmartExplain_AI_executed.ipynb
    └── core/
        ├── __init__.py
        ├── model.py               # LinearRegressionGD
        ├── optimizers.py          # Momentum, LR decay
        ├── feature_engineering.py # California
        └── feature_engineering_house_prices.py # Kaggle

```

```
|- metrics.py          # MAE, MSE, RMSE, R2 (manual)
  └── explainability.py # Feature contributions
  visualization/
    ├── plots.py
    └── cost_surface.py
  simulator/
    └── what_if.py
  app/
    └── streamlit_app.py
  train.py
  model.pkl
  requirements.txt
  README.md
  PROJECT_DETAILS.md
```

## 14. Results

### 14.1 Kaggle House Prices (train.csv)

Metric	Train 80%	Test	Full Data
R <sup>2</sup>	0.87	<b>0.90</b>	0.88
MAE	\$18,821	\$16,410	—
RMSE	\$28,914	\$22,751	—

### 14.2 California Housing (housing.csv)

Metric	Typical Value
R <sup>2</sup>	~0.67–0.70
MAE	~48,000–53,000
RMSE	~66,000–73,000

## Quick Commands

```
# Install
pip install -r requirements.txt

# Train
python train.py

# Run app
streamlit run app/streamlit_app.py

# Run notebook
jupyter notebook notebooks/SmartExplain_AI.ipynb
```

```
# or
python -m jupyter nbconvert --to notebook --execute
notebooks/SmartExplain_AI.ipynb --output SmartExplain_AI_executed.ipynb
```